


Assignment 2 - Spirals

Caleb Zulawski

The implemented perceptron takes 2 inputs (x , y) and returns two outputs ($p(t=0)$, $p(t=1)$). It has 2 hidden layers with 20 neurons each. It is able to classify with near perfect accuracy; classifying the data points near the center of the spiral is nearly impossible since the noise hides any information about which arm the data is from.

Spiral data predictions (accuracy: 0.985)



```

1  #!/usr/bin/env python3
2
3  # Caleb Zulawski
4  # ECE411 - Computational Graphs for Machine Learning
5  # Assignment 2 - Binary classification on a spiral dataset
6
7  import tensorflow as tf
8  import numpy as np
9  import matplotlib.pyplot as plt
10
11 def get_data(N):
12     a = 2;
13     t = np.random.uniform(-3, 3, N)
14     c = t > 0
15     sign = np.where(c, 1, -1)
16     x = a*t*np.cos(np.pi*t)
17     y = a*t*np.sin(np.pi*t) * sign
18     x += np.random.normal(0, 0.05, N);
19     y += np.random.normal(0, 0.05, N);
20     return x, y, c;
21
22 class Model():
23     def __init__(self, sess, n_epochs, learning_rate):
24         self.sess = sess
25         self.n_epochs = n_epochs
26         self.learning_rate = learning_rate
27         self.build_model()
28
29     def build_model(self):
30         n_hidden_1 = 20
31         n_hidden_2 = 20
32         n_input = 2
33         n_output = 2
34
35         self.x = tf.placeholder(tf.float32, [n_input, 1])
36         self.y = tf.placeholder(tf.float32, [1, n_output])
37
38         w1 = tf.get_variable(shape=[n_input, n_hidden_1], dtype=tf.float32, name="weight1", initializer=tf.random_norm
al_initializer())
39         w2 = tf.get_variable(shape=[n_hidden_1, n_hidden_2], dtype=tf.float32, name="weight2", initializer=tf.random_n
ormal_initializer())
40         wo = tf.get_variable(shape=[n_hidden_2, n_output], dtype=tf.float32, name="weighto", initializer=tf.random_nor
mal_initializer())
41
42         b1 = tf.get_variable(shape=[n_hidden_1], dtype=tf.float32, name="bias1", initializer=tf.random_normal_initiali
zer())
43         b2 = tf.get_variable(shape=[n_hidden_2], dtype=tf.float32, name="bias2", initializer=tf.random_normal_initiali
zer())
44         bo = tf.get_variable(shape=[n_output], dtype=tf.float32, name="biaso", initializer=tf.random_normal_initialize
r())
45

```

```

46     layer_1 = tf.add(tf.matmul(tf.transpose(self.x), w1), b1)
47     layer_1 = tf.nn.relu(layer_1)
48     layer_2 = tf.add(tf.matmul(layer_1, w2), b2)
49     layer_2 = tf.nn.relu(layer_2)
50
51     self.layer_out = tf.matmul(layer_2, wo) + bo
52     self.out = tf.nn.softmax(self.layer_out)
53
54     def train(self, xs, ys, cs):
55         cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=self.layer_out, labels=self.y))
56         vs = tf.trainable_variables()
57         cost += tf.add_n([ tf.nn.l2_loss(v) for v in vs if 'bias' not in v.name ]) * 0.001
58         optimizer = tf.train.AdamOptimizer(learning_rate=self.learning_rate).minimize(cost)
59         self.sess.run(tf.global_variables_initializer())
60         for epoch in range(self.n_epochs):
61             for x, y, c in zip(xs, ys, cs):
62                 v = np.expand_dims(np.asarray([x, y]), axis=1)
63                 l = np.reshape([not(c).astype(float), c.astype(float)], (1, 2))
64                 self.sess.run(optimizer, feed_dict={self.x: v, self.y: l})
65
66     def predict(self, x, y):
67         v = np.expand_dims(np.asarray([x, y]), axis=1)
68         return self.sess.run(self.out, feed_dict={self.x: v})[0][1] > 0.5
69
70 n_train = 1000; # if the training set is sparser, sometimes there are gaps in the spiral and it doesn't train well
71 n_test = 1000;
72 x_train, y_train, c_train = get_data(n_train)
73 x_test, y_test, c_test = get_data(n_test)
74
75 with tf.Session() as sess:
76     model = Model(sess, 200, 0.01)
77     model.train(x_train, y_train, c_train)
78
79     total = 0;
80     pred = np.array([]);
81     for x, y, c in zip(x_test, y_test, c_test):
82         pred = np.append(pred, model.predict(x, y))
83
84     plt.figure()
85     plt.xlabel('x')
86     plt.ylabel('y', rotation=0)
87     plt.title('Spiral data predictions (accuracy: {})'.format(np.sum(np.equal(c_test, pred))/n_test))
88     plt.plot(x_test[pred == True], y_test[pred == True], 'ro')
89     plt.plot(x_test[pred == False], y_test[pred == False], 'bo')
90     plt.show()

```