

Caleb Zulawski

## Homework 3 - MNIST

My convolutional network network obtained slightly better than 98.0% accuracy.

## Homework 4 - CIFAR-10/100

State of the art performance for the CIFAR-10 dataset seems to be 95+%, and about 75% for CIFAR-100. I implemented a residual network which scored about 74.4% on CIFAR-10 and 33.9% on CIFAR-100. This is not very close to current performance, but was not unexpected due to the network construction. To reduce training time on available hardware I used only 4 residual blocks, as opposed to the 16 in ResNet-50, and even more in larger networks. I also had to reduce the number of features, and train on the original 32x32 images rather than upsampled 224x224 that many networks seem to use. It's also possible that the values used were not optimal for a 4 residual block network.

Beyond increasing the size of the network, I would have liked to try more optimizers. I used the momentum optimizer with a momentum of 0.9, based on some successes I saw in other papers. I would have also liked to use a learning rate schedule to speed up early training while training better in later epochs, but I compromised and used a constant rate of 0.01. Finally, I would have liked to use augmented data.

```

1  #!/usr/bin/env python3
2
3  import tensorflow as tf
4  import tensorflow.contrib.slim as slim
5  from tensorflow.contrib.learn.python.learn.datasets import mnist
6  import numpy as np
7
8  class Model():
9      def __init__(self, sess, pixels, n_batch, n_classes, learning_rate):
10         self.sess = sess
11         self.learning_rate = learning_rate
12         self.pixels = pixels
13         self.n_batch = n_batch
14         self.n_classes = n_classes
15         self.build_model()
16
17     def build_model(self):
18         self.inputs = tf.placeholder(tf.float32, [None, self.pixels, self.pixels, 1])
19         self.labels = tf.placeholder(tf.float32, [None, self.n_classes])
20         with slim.arg_scope([slim.conv2d, slim.fully_connected],
21                             activation_fn=tf.nn.relu,
22                             weights_initializer=tf.truncated_normal_initializer(0.0, 0.01),
23                             weights_regularizer=slim.l2_regularizer(0.0005)):
24             net = self.inputs
25             net = slim.repeat(net, 1, slim.conv2d, 32, [5, 5], scope='conv1')
26             net = slim.max_pool2d(net, [2, 2], scope='pool1')
27             net = slim.repeat(net, 1, slim.conv2d, 64, [5, 5], scope='conv2')
28             net = slim.max_pool2d(net, [2, 2], scope='pool2')
29             net = slim.flatten(net, scope='flatten2')
30             net = slim.fully_connected(net, 1024, scope='fc3')
31             net = slim.dropout(net, 0.8, scope='dropout3')
32             net = slim.fully_connected(net, self.n_classes, activation_fn=None, normalizer_fn=None, scope='fc4')
33         self.predictions = net
34         slim.losses.softmax_cross_entropy(self.predictions, self.labels)
35         self.loss = slim.losses.get_total_loss(add_regularization_losses=True)
36         self.optimizer = tf.train.AdamOptimizer(learning_rate=self.learning_rate).minimize(self.loss)
37
38         correct = tf.equal(tf.argmax(self.predictions, 1), tf.argmax(self.labels, 1))
39         self.accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
40
41     def train_minibatch(self, xs, ys):
42         self.sess.run(self.optimizer, feed_dict={self.inputs: xs, self.labels: ys})
43
44     def train(self, xs, ys, xvs, yvs, epochs):
45         num_minibatches = xs.shape[0] // self.n_batch
46         for epoch in range(epochs):
47             p = np.random.permutation(xs.shape[0])
48             for i in range(num_minibatches):
49                 # print("Minibatch " + str(i) + '/' + str(num_minibatches))
50                 start = i * self.n_batch
51                 end = (i + 1) * self.n_batch

```

```
52         self.train_minibatch(xs[p][start:end], ys[p][start:end])
53         accuracy = self.validate(xvs, yvs)
54         print("Epoch {} validation accuracy: {}".format(epoch, accuracy))
55
56     def validate(self, xs, ys):
57         return self.sess.run(self.accuracy, feed_dict={self.inputs: xs, self.labels: ys})
58
59     def test(self, xs):
60         return self.sess.run(self.predictions, feed_dict={self.inputs: xs})
61
62 with tf.Session() as sess:
63     data = mnist.read_data_sets("MNIST_data/", dtype=tf.uint8, reshape=False, one_hot=True)
64     m = Model(sess, 28, 50, 10, 0.001)
65     sess.run(tf.global_variables_initializer())
66     m.train(data.train.images, data.train.labels, data.validation.images, data.validation.labels, 2)
67     test_accuracy = m.validate(data.test.images, data.test.labels)
68     print("Test set accuracy: {}".format(test_accuracy))
```

```

1  #!/usr/bin/env python3
2
3  import tensorflow as tf
4  import tensorflow.contrib.slim as slim
5  import numpy as np
6
7  # from tflearn.datasets import cifar10
8  from tflearn.datasets import cifar100
9
10 def residual(inputs, depth, kernel, scope='residual'):
11     with tf.variable_scope(scope, 'residual', [inputs]) as sc:
12         residual = inputs
13         residual = slim.batch_norm(residual, activation_fn=tf.nn.relu, scope='bn1')
14         residual = slim.conv2d(residual, depth, kernel, scope='conv1')
15         residual = slim.batch_norm(residual, activation_fn=tf.nn.relu, scope='bn2')
16         residual = slim.conv2d(residual, inputs.get_shape()[3], kernel, scope='conv2')
17         residual = residual + inputs
18     return residual
19
20 def resnet_small(inputs,
21                 num_classes=None,
22                 global_pool=True,
23                 output_stride=None,
24                 reuse=None,
25                 scope='resnet_small'):
26     with tf.variable_scope(scope, 'resnet_small', [inputs]) as sc:
27         net = inputs
28         net = slim.repeat(net, 1, slim.conv2d, 32, [5, 5], stride=1, scope='conv1')
29         net = slim.max_pool2d(net, [2, 2], scope='pool1')
30         net = slim.repeat(net, 1, slim.conv2d, 32, [5, 5], stride=1, scope='conv2')
31         net = slim.max_pool2d(net, [2, 2], scope='pool2')
32         net = residual(net, 32, [3, 3], scope='block3')
33         net = residual(net, 32, [3, 3], scope='block4')
34         net = residual(net, 64, [3, 3], scope='block5')
35         net = residual(net, 64, [3, 3], scope='block6')
36         net = slim.flatten(net, scope='flatten6')
37         net = slim.fully_connected(net, 1024, scope='fc7')
38         net = slim.dropout(net, 0.5, scope='dropout8')
39         net = slim.fully_connected(net, num_classes, activation_fn=None, normalizer_fn=None, scope='logits')
40     return net
41
42 class Model():
43     def __init__(self, sess, n_batch, n_classes, learning_rate):
44         self.sess = sess
45         self.learning_rate = learning_rate
46         self.n_batch = n_batch
47         self.n_classes = n_classes
48         self.build_model()
49
50     def build_model(self):
51         self.inputs = tf.placeholder(tf.float32, [None, 32, 32, 3])

```

```

52     self.labels = tf.placeholder(tf.float32, [None, self.n_classes])
53     with slim.arg_scope([slim.conv2d, slim.fully_connected],
54                         activation_fn=tf.nn.relu,
55                         weights_initializer=tf.truncated_normal_initializer(0.0, 0.01),
56                         weights_regularizer=slim.l2_regularizer(0.002)):
57         net = self.inputs
58         net = resnet_small(net, num_classes=self.n_classes)
59         self.predictions = net
60         slim.losses.softmax_cross_entropy(self.predictions, self.labels)
61         self.loss = slim.losses.get_total_loss(add_regularization_losses=True)
62         self.optimizer = tf.train.MomentumOptimizer(learning_rate=self.learning_rate, momentum=0.9).minimize(self.loss)
63
64     correct = tf.equal(tf.argmax(self.predictions, 1), tf.argmax(self.labels, 1))
65     self.accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
66
67     def train_minibatch(self, xs, ys):
68         self.sess.run(self.optimizer, feed_dict={self.inputs: xs, self.labels: ys})
69
70     def train(self, xs, ys, xvs, yvs, epochs):
71         num_minibatches = xs.shape[0] // self.n_batch
72         for epoch in range(epochs):
73             p = np.random.permutation(xs.shape[0])
74             for i in range(num_minibatches):
75                 # print("Minibatch " + str(i) + '/' + str(num_minibatches))
76                 start = i * self.n_batch
77                 end = (i + 1) * self.n_batch
78                 self.train_minibatch(xs[p][start:end], ys[p][start:end])
79             accuracy = self.validate(xvs, yvs)
80             print("Epoch {} validation accuracy: {}".format(epoch, accuracy))
81
82     def validate(self, xs, ys):
83         return self.sess.run(self.accuracy, feed_dict={self.inputs: xs, self.labels: ys})
84
85     def test(self, xs):
86         return self.sess.run(self.predictions, feed_dict={self.inputs: xs})
87
88     with tf.Session() as sess:
89         # (train_images, train_labels), (test_images, test_labels) = cifar10.load_data(one_hot=True)
90         # m = Model(sess, 128, 10, 0.01)
91         (train_images, train_labels), (test_images, test_labels) = cifar100.load_data(one_hot=True)
92         m = Model(sess, 128, 100, 0.01)
93         sess.run(tf.global_variables_initializer())
94         validation_images = train_images[0:5000]
95         validation_labels = train_labels[0:5000]
96         train_images = train_images[5000:]
97         train_labels = train_labels[5000:]
98         m.train(train_images, train_labels, validation_images, validation_labels, 100)
99         test_accuracy = m.validate(test_images, test_labels)
100        print("Test set accuracy: {}".format(test_accuracy))

```