

```

1  #!/usr/bin/env python3
2
3  import tensorflow as tf
4  import tensorflow.contrib.slim as slim
5  import numpy as np
6
7  # from tflearn.datasets import cifar10
8  from tflearn.datasets import cifar100
9
10 def residual(inputs, depth, kernel, scope='residual'):
11     with tf.variable_scope(scope, 'residual', [inputs]) as sc:
12         residual = inputs
13         residual = slim.batch_norm(residual, activation_fn=tf.nn.relu, scope='bn1')
14         residual = slim.conv2d(residual, depth, kernel, scope='conv1')
15         residual = slim.batch_norm(residual, activation_fn=tf.nn.relu, scope='bn2')
16         residual = slim.conv2d(residual, inputs.get_shape()[3], kernel, scope='conv2')
17         residual = residual + inputs
18     return residual
19
20 def resnet_small(inputs,
21                 num_classes=None,
22                 global_pool=True,
23                 output_stride=None,
24                 reuse=None,
25                 scope='resnet_small'):
26     with tf.variable_scope(scope, 'resnet_small', [inputs]) as sc:
27         net = inputs
28         net = slim.repeat(net, 1, slim.conv2d, 32, [5, 5], stride=1, scope='conv1')
29         net = slim.max_pool2d(net, [2, 2], scope='pool1')
30         net = slim.repeat(net, 1, slim.conv2d, 32, [5, 5], stride=1, scope='conv2')
31         net = slim.max_pool2d(net, [2, 2], scope='pool2')
32         net = residual(net, 32, [3, 3], scope='block3')
33         net = residual(net, 32, [3, 3], scope='block4')
34         net = residual(net, 64, [3, 3], scope='block5')
35         net = residual(net, 64, [3, 3], scope='block6')
36         net = slim.flatten(net, scope='flatten6')
37         net = slim.fully_connected(net, 1024, scope='fc7')
38         net = slim.dropout(net, 0.5, scope='dropout8')
39         net = slim.fully_connected(net, num_classes, activation_fn=None, normalizer_fn=None, scope='logits')
40     return net
41
42 class Model():
43     def __init__(self, sess, n_batch, n_classes, learning_rate):
44         self.sess = sess
45         self.learning_rate = learning_rate
46         self.n_batch = n_batch
47         self.n_classes = n_classes
48         self.build_model()
49
50     def build_model(self):
51         self.inputs = tf.placeholder(tf.float32, [None, 32, 32, 3])

```

```

52     self.labels = tf.placeholder(tf.float32, [None, self.n_classes])
53     with slim.arg_scope([slim.conv2d, slim.fully_connected],
54                         activation_fn=tf.nn.relu,
55                         weights_initializer=tf.truncated_normal_initializer(0.0, 0.01),
56                         weights_regularizer=slim.l2_regularizer(0.002)):
57         net = self.inputs
58         net = resnet_small(net, num_classes=self.n_classes)
59     self.predictions = net
60     slim.losses.softmax_cross_entropy(self.predictions, self.labels)
61     self.loss = slim.losses.get_total_loss(add_regularization_losses=True)
62     self.optimizer = tf.train.MomentumOptimizer(learning_rate=self.learning_rate, momentum=0.9).minimize(self.loss)
63
64     correct = tf.equal(tf.argmax(self.predictions, 1), tf.argmax(self.labels, 1))
65     self.accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
66
67     def train_minibatch(self, xs, ys):
68         self.sess.run(self.optimizer, feed_dict={self.inputs: xs, self.labels: ys})
69
70     def train(self, xs, ys, xvs, yvs, epochs):
71         num_minibatches = xs.shape[0] // self.n_batch
72         for epoch in range(epochs):
73             p = np.random.permutation(xs.shape[0])
74             for i in range(num_minibatches):
75                 # print("Minibatch " + str(i) + '/' + str(num_minibatches))
76                 start = i * self.n_batch
77                 end = (i + 1) * self.n_batch
78                 self.train_minibatch(xs[p][start:end], ys[p][start:end])
79             accuracy = self.validate(xvs, yvs)
80             print("Epoch {} validation accuracy: {}".format(epoch, accuracy))
81
82     def validate(self, xs, ys):
83         return self.sess.run(self.accuracy, feed_dict={self.inputs: xs, self.labels: ys})
84
85     def test(self, xs):
86         return self.sess.run(self.predictions, feed_dict={self.inputs: xs})
87
88     with tf.Session() as sess:
89         # (train_images, train_labels), (test_images, test_labels) = cifar10.load_data(one_hot=True)
90         # m = Model(sess, 128, 10, 0.01)
91         (train_images, train_labels), (test_images, test_labels) = cifar100.load_data(one_hot=True)
92         m = Model(sess, 128, 100, 0.01)
93         sess.run(tf.global_variables_initializer())
94         validation_images = train_images[0:5000]
95         validation_labels = train_labels[0:5000]
96         train_images = train_images[5000:]
97         train_labels = train_labels[5000:]
98         m.train(train_images, train_labels, validation_images, validation_labels, 100)
99         test_accuracy = m.validate(test_images, test_labels)
100        print("Test set accuracy: {}".format(test_accuracy))

```